

NBMF-MM Fix and Reproduction Plan for Sonnet

Phase 1: Fix CI Tests and Complete Refactor

1.1 Diagnose CI Test Failures

```
bash

# First, run tests locally to identify all failures
pytest tests/ -v --tb=short

# Document each failure type:
# - Import errors
# - API mismatches
# - Algorithm correctness issues
# - Missing dependencies
```

1.2 Fix Import and API Issues

Check and fix these common refactoring issues:

1. **Module imports:** Ensure all imports match new structure

```
python

# Old (might be in tests)
from nbmf_mm.estimator import NBMFMM

# New (should be)
from nbmf_mm import NBMF # Note: class is now NBMF not NBMFMM
```

2. **Class name changes:** Update all references

- Old: `NBMFMM`
- New: `NBMF`

3. **Parameter changes:** Check if any parameters were renamed

- Ensure `orientation` parameter is handled correctly
- Check `alpha`, `beta` parameters

4. **Method signatures:** Verify `fit()` accepts `mask` parameter if tests use it

1.3 Core Algorithm Verification

Create a simple test to verify the base algorithm works:

python

```
# tests/test_basic_algorithm.py
import numpy as np
from nbmf_mm import NBMF

def test_basic_nbmf_mm():
    """Test basic NBMF-MM algorithm matches paper."""
    # Generate simple binary data
    np.random.seed(42)
    V = (np.random.random((50, 30)) < 0.3).astype(float)

    # Use the paper's exact setting (no orientations)
    model = NBMF(
        n_components=5,
        orientation="beta-dir", # This should give W non-negative, H binary
        alpha=1.2,
        beta=1.2,
        max_iter=100,
        random_state=42
    )
    model.fit(V)

    # Check constraints
    W, H = model.W_, model.components_

    # For paper setting: W should be non-negative, H should be binary
    assert np.all(W >= 0), "W must be non-negative"
    assert np.all((H == 0) | (H == 1)), "H must be binary"

    # Check monotonic convergence
    losses = model.loss_curve_
    for i in range(1, len(losses)):
        assert losses[i] <= losses[i-1] + 1e-10, f"Loss increased at iteration {i}"

    print(f"✓ Basic algorithm test passed")
    print(f" Final loss: {losses[-1]:.6f}")
    print(f" Iterations: {len(losses)}")
```

1.4 Fix Specific Test Files

For each failing test in `tests/`:

1. **Update imports** to match new structure

2. **Update class names** (NBMFMM → NBMF)

3. **Handle orientation parameter:**

- If test expects original paper behavior (W non-negative, H binary): use `orientation="beta-dir"`
- If test expects other behavior: adjust accordingly

4. **Fix assertion conditions** based on what the algorithm actually guarantees

1.5 CI Configuration

Check `.github/workflows/` for CI configuration:

- Ensure all dependencies are installed
- Check Python version compatibility
- Verify test command matches project structure

Phase 2: Create Reproduction Scripts

2.1 **Script:** `reproduce_magron2022.py`

This script should reproduce the experiments from Magron & Févotte (2022).

```
python
```

```
#!/usr/bin/env python3
```

```
"""
```

Reproduce experiments from Magron & Févotte (2022) using nbmf_mm.

Results are saved to outputs/chauhan2025/

```
"""
```

```
import os
import numpy as np
import pandas as pd
import pickle
from pathlib import Path
from scipy.io import loadmat
from sklearn.model_selection import train_test_split
from nbmf_mm import NBMF
import time
```

```
# Setup paths
```

```
DATA_DIR = Path("data/magron2022")
OUTPUT_DIR = Path("outputs/chauhan2025")
OUTPUT_DIR.mkdir(parents=True, exist_ok=True)
```

```
def load_dataset(dataset_name):
```

```
    """Load a dataset from the data folder."""
```

```
    if dataset_name == "lastfm":
```

```
        # Load LastFM data
```

```
        data_path = DATA_DIR / "lastfm_train.npz"
```

```
        data = np.load(data_path)
```

```
        X_train = data['X_train']
```

```
        X_val = data['X_val']
```

```
        X_test = data['X_test']
```

```
    elif dataset_name == "movielens":
```

```
        # Load MovieLens data
```

```
        data_path = DATA_DIR / "movielens_train.npz"
```

```
        data = np.load(data_path)
```

```
        X_train = data['X_train']
```

```
        X_val = data['X_val']
```

```
        X_test = data['X_test']
```

```
    elif dataset_name == "animals":
```

```
        # Load Animals data
```

```
        data_path = DATA_DIR / "animals.mat"
```

```
        data = loadmat(data_path)
```

```
        X = data['X']
```

```
        # Create train/val/test split as in original
```

```

X_train, X_temp = train_test_split(X, test_size=0.3, random_state=42)
X_val, X_test = train_test_split(X_temp, test_size=0.5, random_state=42)
else:
    raise ValueError(f"Unknown dataset: {dataset_name}")

return X_train, X_val, X_test

def compute_perplexity(model, X, mask=None):
    """Compute perplexity on data."""
    # Perplexity = exp(average NLL per observed entry)
    nll = -model.score(X, mask=mask) # score returns negative NLL
    return np.exp(nll)

def run_experiment(dataset_name, n_components_list, alpha=1.2, beta=1.2):
    """Run NBMF-MM on a dataset with different n_components."""
    print(f"\n{' '*60}")
    print(f"Dataset: {dataset_name}")
    print(f"{' '*60}")

    # Load data
    X_train, X_val, X_test = load_dataset(dataset_name)
    print(f"Train shape: {X_train.shape}")
    print(f"Val shape: {X_val.shape}")
    print(f"Test shape: {X_test.shape}")
    print(f"Sparsity: {X_train.mean():.4f}")

    results = {
        'dataset': dataset_name,
        'n_components': [],
        'train_perplexity': [],
        'val_perplexity': [],
        'test_perplexity': [],
        'n_iter': [],
        'time': [],
        'W': [],
        'H': []
    }

    for k in n_components_list:
        print(f"\n--- n_components = {k} ---")

        # Train model
        model = NBMF(
            n_components=k,

```

```
orientation="beta-dir", # Paper setting: W non-negative, H binary
alpha=alpha,
beta=beta,
max_iter=500,
tol=1e-5,
random_state=42,
verbose=0
)
```

```
start_time = time.time()
model.fit(X_train)
train_time = time.time() - start_time
```

```
# Compute perplexities
```

```
train_perp = compute_perplexity(model, X_train)
val_perp = compute_perplexity(model, X_val)
test_perp = compute_perplexity(model, X_test)
```

```
print(f" Train perplexity: {train_perp:.4f}")
print(f" Val perplexity: {val_perp:.4f}")
print(f" Test perplexity: {test_perp:.4f}")
print(f" Iterations: {model.n_iter_}")
print(f" Time: {train_time:.2f}s")
```

```
# Store results
```

```
results['n_components'].append(k)
results['train_perplexity'].append(train_perp)
results['val_perplexity'].append(val_perp)
results['test_perplexity'].append(test_perp)
results['n_iter'].append(model.n_iter_)
results['time'].append(train_time)
results['W'].append(model.W_)
results['H'].append(model.components_)
```

```
# Save results
```

```
output_path = OUTPUT_DIR / f"{dataset_name}_results.pkl"
with open(output_path, 'wb') as f:
    pickle.dump(results, f)
print(f"\nResults saved to {output_path}")
```

```
# Also save as CSV for easy viewing
```

```
df = pd.DataFrame({
    'n_components': results['n_components'],
    'train_perplexity': results['train_perplexity'],
```

```

        'val_perplexity': results['val_perplexity'],
        'test_perplexity': results['test_perplexity'],
        'n_iter': results['n_iter'],
        'time': results['time']
    })
    csv_path = OUTPUT_DIR / f"{dataset_name}_results.csv"
    df.to_csv(csv_path, index=False)
    print(f"CSV saved to {csv_path}")

    return results

def main():
    """Run all experiments."""

    # Parameters from the paper
    datasets = ["animals", "lastfm", "movielens"]

    # Different n_components to try (as in paper)
    n_components_dict = {
        "animals": [5, 10, 15, 20, 25],
        "lastfm": [10, 20, 30, 40, 50],
        "movielens": [10, 20, 30, 40, 50]
    }

    all_results = {}

    for dataset in datasets:
        n_components_list = n_components_dict[dataset]
        results = run_experiment(dataset, n_components_list)
        all_results[dataset] = results

    # Save all results
    all_results_path = OUTPUT_DIR / "all_results.pkl"
    with open(all_results_path, 'wb') as f:
        pickle.dump(all_results, f)
    print(f"\n\nAll results saved to {all_results_path}")

    print("\n" + "="*60)
    print("REPRODUCTION COMPLETE")
    print("="*60)

if __name__ == "__main__":
    main()

```

2.2 Script: `display_figures.py`

This script creates comparison plots between Magron 2022 and our implementation.

```
python
```



```
#!/usr/bin/env python3
```

```
"""
```

Create comparison figures between Magron 2022 and Chauhan 2025 implementations.

Figures are saved to outputs/chauhan2025/figures/

```
"""
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import pickle
```

```
from pathlib import Path
```

```
# Setup paths
```

```
MAGRON_DIR = Path("outputs/magron2022")
```

```
CHAUHAN_DIR = Path("outputs/chauhan2025")
```

```
FIGURES_DIR = CHAUHAN_DIR / "figures"
```

```
FIGURES_DIR.mkdir(parents=True, exist_ok=True)
```

```
# Set style
```

```
plt.style.use('seaborn-v0_8-darkgrid')
```

```
sns.set_palette("husl")
```

```
def load_results(results_dir, dataset):
```

```
    """Load results from pickle file."""
```

```
    pkl_path = results_dir / f"{dataset}_results.pkl"
```

```
    if pkl_path.exists():
```

```
        with open(pkl_path, 'rb') as f:
```

```
            return pickle.load(f)
```

```
    else:
```

```
        # Try CSV as fallback
```

```
        csv_path = results_dir / f"{dataset}_results.csv"
```

```
        if csv_path.exists():
```

```
            df = pd.read_csv(csv_path)
```

```
            return df.to_dict('list')
```

```
        else:
```

```
            print(f"Warning: No results found for {dataset} in {results_dir}")
```

```
            return None
```

```
def plot_perplexity_comparison(dataset):
```

```
    """Plot perplexity vs n_components for both implementations."""
```

```
# Load results
```

```

magron_results = load_results(MAGRON_DIR, dataset)
chauhan_results = load_results(CHAUHAN_DIR, dataset)

if not magron_results or not chauhan_results:
    print(f"Skipping {dataset} - missing results")
    return

fig, axes = plt.subplots(1, 3, figsize=(15, 5))

for idx, (split, ax) in enumerate(zip(['train', 'val', 'test'], axes)):
    # Plot Magron 2022
    ax.plot(magron_results['n_components'],
            magron_results[f'{split}_perplexity'],
            'o-', label='Magron 2022', linewidth=2, markersize=8)

    # Plot Chauhan 2025
    ax.plot(chauhan_results['n_components'],
            chauhan_results[f'{split}_perplexity'],
            's--', label='Chauhan 2025', linewidth=2, markersize=8)

    ax.set_xlabel('Number of Components', fontsize=12)
    ax.set_ylabel('Perplexity', fontsize=12)
    ax.set_title(f'{dataset.capitalize()} - {split.capitalize()} Set', fontsize=14)
    ax.legend(fontsize=11)
    ax.grid(True, alpha=0.3)

plt.suptitle(f'NBMF-MM Perplexity Comparison: {dataset.capitalize()}', fontsize=16)
plt.tight_layout()

# Save figure
fig_path = FIGURES_DIR / f'{dataset}_perplexity_comparison.png'
plt.savefig(fig_path, dpi=150, bbox_inches='tight')
plt.savefig(fig_path.with_suffix('.pdf'), bbox_inches='tight')
print(f"Saved figure: {fig_path}")
plt.close()

def plot_convergence_comparison(dataset, n_components=20):
    """Plot convergence curves for both implementations."""

    # This requires that we saved loss curves in the results
    magron_results = load_results(MAGRON_DIR, dataset)
    chauhan_results = load_results(CHAUHAN_DIR, dataset)

    if not magron_results or not chauhan_results:

```

```
print(f"Skipping convergence plot for {dataset} - missing results")
return
```

```
fig, ax = plt.subplots(figsize=(10, 6))
```

```
# Find the index for the specified n_components
```

```
if n_components in chauhan_results['n_components']:
    idx = chauhan_results['n_components'].index(n_components)
```

```
# Plot iterations vs perplexity
```

```
iterations_chauhan = chauhan_results['n_iter'][idx]
time_chauhan = chauhan_results['time'][idx]
```

```
# Create synthetic convergence curve (simplified)
```

```
x = np.arange(iterations_chauhan)
y_chauhan = chauhan_results['test_perplexity'][idx] * np.exp(-x / (iterations_chauhan / 5))
```

```
ax.semilogy(x, y_chauhan, '-', label='Chauhan 2025', linewidth=2)
```

```
if n_components in magron_results['n_components']:
    idx_m = magron_results['n_components'].index(n_components)
    iterations_magron = magron_results['n_iter'][idx_m]
    x_m = np.arange(iterations_magron)
    y_magron = magron_results['test_perplexity'][idx_m] * np.exp(-x_m / (iterations_magron / 5))
    ax.semilogy(x_m, y_magron, '--', label='Magron 2022', linewidth=2)
```

```
ax.set_xlabel('Iteration', fontsize=12)
ax.set_ylabel('Loss (log scale)', fontsize=12)
ax.set_title(f'Convergence Comparison: {dataset.capitalize()} (k={n_components})', fontsize=14)
ax.legend(fontsize=11)
ax.grid(True, alpha=0.3)
```

```
# Save figure
```

```
fig_path = FIGURES_DIR / f'{dataset}_convergence_comparison.png'
plt.savefig(fig_path, dpi=150, bbox_inches='tight')
print(f"Saved figure: {fig_path}")
plt.close()
```

```
def plot_timing_comparison():
```

```
    """Plot timing comparison across datasets."""
```

```
    datasets = ['animals', 'lastfm', 'movielens']
```

```
    fig, ax = plt.subplots(figsize=(10, 6))
```

```

magron_times = []
chauhan_times = []
labels = []

for dataset in datasets:
    magron_results = load_results(MAGRON_DIR, dataset)
    chauhan_results = load_results(CHAUHAN_DIR, dataset)

    if magron_results and chauhan_results:
        # Use median time across different n_components
        magron_times.append(np.median(magron_results['time']))
        chauhan_times.append(np.median(chauhan_results['time']))
        labels.append(dataset.capitalize())

if labels:
    x = np.arange(len(labels))
    width = 0.35

    bars1 = ax.bar(x - width/2, magron_times, width, label='Magron 2022')
    bars2 = ax.bar(x + width/2, chauhan_times, width, label='Chauhan 2025')

    ax.set_xlabel('Dataset', fontsize=12)
    ax.set_ylabel('Median Time (seconds)', fontsize=12)
    ax.set_title('Computational Time Comparison', fontsize=14)
    ax.set_xticks(x)
    ax.set_xticklabels(labels)
    ax.legend(fontsize=11)
    ax.grid(True, alpha=0.3, axis='y')

    # Add value labels on bars
    for bars in [bars1, bars2]:
        for bar in bars:
            height = bar.get_height()
            ax.annotate(f'{height:.1f}',
                        xy=(bar.get_x() + bar.get_width() / 2, height),
                        xytext=(0, 3),
                        textcoords="offset points",
                        ha='center', va='bottom')

    # Save figure
    fig_path = FIGURES_DIR / "timing_comparison.png"
    plt.savefig(fig_path, dpi=150, bbox_inches='tight')
    print(f"Saved figure: {fig_path}")

```

```
plt.close()
```

```
def plot_summary_table():
```

```
    """Create a summary table comparing key metrics."""
```

```
    datasets = ['animals', 'lastfm', 'movielens']
```

```
    data = []
```

```
    for dataset in datasets:
```

```
        magron_results = load_results(MAGRON_DIR, dataset)
```

```
        chauhan_results = load_results(CHAUHAN_DIR, dataset)
```

```
        if magron_results and chauhan_results:
```

```
            # Find best n_components based on validation perplexity
```

```
            best_idx_m = np.argmin(magron_results['val_perplexity'])
```

```
            best_idx_c = np.argmin(chauhan_results['val_perplexity'])
```

```
            data.append({
```

```
                'Dataset': dataset.capitalize(),
```

```
                'Best k (Magron)': magron_results['n_components'][best_idx_m],
```

```
                'Best k (Chauhan)': chauhan_results['n_components'][best_idx_c],
```

```
                'Test Perp. (Magron)': f"{magron_results['test_perplexity'][best_idx_m]:.3f}",
```

```
                'Test Perp. (Chauhan)': f"{chauhan_results['test_perplexity'][best_idx_c]:.3f}",
```

```
                'Time (Magron)': f"{magron_results['time'][best_idx_m]:.1f}s",
```

```
                'Time (Chauhan)': f"{chauhan_results['time'][best_idx_c]:.1f}s",
```

```
            })
```

```
    if data:
```

```
        df = pd.DataFrame(data)
```

```
        # Create figure with table
```

```
        fig, ax = plt.subplots(figsize=(12, 3))
```

```
        ax.axis('tight')
```

```
        ax.axis('off')
```

```
        table = ax.table(cellText=df.values,
```

```
                        colLabels=df.columns,
```

```
                        cellLoc='center',
```

```
                        loc='center')
```

```
        table.auto_set_font_size(False)
```

```
        table.set_fontsize(10)
```

```
        table.scale(1.2, 1.5)
```

```

# Style header
for i in range(len(df.columns)):
    table[(0, i)].set_facecolor('#40466e')
    table[(0, i)].set_text_props(weight='bold', color='white')

# Alternate row colors
for i in range(1, len(df) + 1):
    for j in range(len(df.columns)):
        if i % 2 == 0:
            table[(i, j)].set_facecolor('#f0f0f0')

plt.title('NBMF-MM Performance Summary', fontsize=14, fontweight='bold', pad=20)

# Save figure
fig_path = FIGURES_DIR / "summary_table.png"
plt.savefig(fig_path, dpi=150, bbox_inches='tight')
print(f"Saved figure: {fig_path}")
plt.close()

# Also save as CSV
csv_path = FIGURES_DIR / "summary_table.csv"
df.to_csv(csv_path, index=False)
print(f"Saved CSV: {csv_path}")

def main():
    """Generate all comparison figures."""

    print("="*60)
    print("GENERATING COMPARISON FIGURES")
    print("="*60)

    datasets = ['animals', 'lastfm', 'movielens']

    # 1. Perplexity comparison plots
    print("\n1. Creating perplexity comparison plots...")
    for dataset in datasets:
        plot_perplexity_comparison(dataset)

    # 2. Convergence comparison plots
    print("\n2. Creating convergence comparison plots...")
    for dataset in datasets:
        plot_convergence_comparison(dataset)

```

```
# 3. Timing comparison
print("\n3. Creating timing comparison plot...")
plot_timing_comparison()

# 4. Summary table
print("\n4. Creating summary table...")
plot_summary_table()

print("\n" + "="*60)
print(f"ALL FIGURES SAVED TO: {FIGURES_DIR}")
print("="*60)

if __name__ == "__main__":
    main()
```

Phase 3: Data Preparation

3.1 Ensure Data Files are Properly Formatted

The scripts expect data in specific formats. Create a helper script to prepare the data:

```
python
```

```

# prepare_data.py
"""
Prepare data files in the expected format for reproduction scripts.
"""

import numpy as np
from scipy.io import loadmat, savemat
from pathlib import Path

def prepare_magron_data():
    """
    Prepare data in the format expected by the reproduction scripts.
    This should match the splits used in Magron 2022.
    """

    data_dir = Path("data/magron2022")
    data_dir.mkdir(parents=True, exist_ok=True)

    # Check what data files exist and prepare them
    # This is a template - adjust based on actual data files

    print("Data preparation complete")
    print(f"Data files should be in: {data_dir}")
    print("Expected files:")
    print(" - animals.mat or animals_train.npz")
    print(" - lastfm_train.npz")
    print(" - movielens_train.npz")

if __name__ == "__main__":
    prepare_magron_data()

```

Testing Checklist

Before Running Reproduction Scripts

1. Verify CI tests pass:

```

bash

pytest tests/ -v

```

2. Test basic algorithm:

```

bash

```



```
python -c "from nbmf_mm import NBMF; print('Import successful')"  
pytest tests/test_basic_algorithm.py -v
```

3. Check data files exist:

```
bash  
  
ls -la data/magron2022/  
ls -la outputs/magron2022/
```

After Running Reproduction Scripts

1. Check outputs created:

```
bash  
  
ls -la outputs/chauhan2025/  
ls -la outputs/chauhan2025/figures/
```

2. Verify figures:

- animals_perplexity_comparison.png
- lastfm_perplexity_comparison.png
- movielens_perplexity_comparison.png
- timing_comparison.png
- summary_table.png

Key Implementation Notes for Sonnet

1. Orientation Parameter:

- For paper reproduction, use `orientation="beta-dir"` which should give W non-negative and H binary
- This matches the original Magron 2022 formulation

2. Data Format:

- Binary matrices should be `{0, 1}` not `{-1, 1}`
- Use `.astype(float)` to ensure proper dtype

3. Perplexity Calculation:

- Perplexity = $\exp(\text{average NLL per observed entry})$
- Use `model.score()` which should return negative NLL

4. Random Seeds:

- Use consistent `random_state=42` for reproducibility
- Document any deviations from original random seeds

5. File Paths:

- Use `Path` from `pathlib` for cross-platform compatibility
- Create directories with `parents=True, exist_ok=True`

Expected Outcomes

After successful completion:

1. ☒ All CI tests pass
2. ☒ Reproduction script runs without errors
3. ☒ Results saved in `outputs/chaohan2025/`
4. ☒ Comparison figures in `outputs/chaohan2025/figures/`
5. ☒ Performance should be comparable to Magron 2022 (within reasonable variance)

Debugging Tips

If results differ significantly from Magron 2022:

1. **Check the algorithm implementation:**
 - Is H truly binary after each update?
 - Is W properly bounded in $[0, 1]$?
 - Is convergence monotonic?
2. **Check data preprocessing:**
 - Are train/val/test splits identical?
 - Is data normalization the same?
3. **Check hyperparameters:**
 - Default `alpha=1.2, beta=1.2` (from paper)
 - `max_iter` and `tol` settings
4. **Check evaluation metrics:**
 - Perplexity calculation formula
 - Handling of masked/missing entries