

NBMF-MM Solver Implementation Specification

Based on Magron & Févotte (2022) Reference Implementation

CRITICAL UNDERSTANDING

The KEY misunderstanding we had:

- **H is NOT binary during optimization!** It's continuous in $[0,1]$
- **W rows sum to 1** (simplex constraint)
- The "binary" in NBMF refers to the **binary input data Y**, not the factors!
- H becomes "effectively binary" at convergence due to the Beta prior, but is NOT forced to $\{0,1\}$ during updates

1. Mathematical Formulation

Objective Function

Minimize negative log-likelihood with Beta prior:

$$L = -\sum[Y * \log(\Theta) + (1-Y) * \log(1-\Theta)] - \sum[(\alpha-1) * \log(H) + (\beta-1) * \log(1-H)]$$

where:

- $Y \in \{0,1\}^{(m \times n)}$ is binary data
- $\Theta = W^T @ H \in (0,1)^{(m \times n)}$ is the Bernoulli parameter
- $W \in \mathbb{R}^{+ (k \times m)}$ with columns summing to 1 (probability simplex)
- $H \in (0,1)^{(k \times n)}$ is continuous (NOT binary during optimization!)
- α, β are Beta prior parameters for H

2. Core Algorithm Implementation

File: `src/nbmf_mm/_solver.py`

```
python
```

```
import numpy as np
import time
from typing import Tuple, Optional, List
```

```
def nbmf_mm_solver(
    Y: np.ndarray,
    n_components: int,
    max_iter: int = 500,
    tol: float = 1e-5,
    alpha: float = 1.2,
    beta: float = 1.2,
    W_init: Optional[np.ndarray] = None,
    H_init: Optional[np.ndarray] = None,
    mask: Optional[np.ndarray] = None,
    random_state: Optional[int] = None,
    verbose: int = 0,
    eps: float = 1e-8
) -> Tuple[np.ndarray, np.ndarray, List[float], float, int]:
    """
```

NBMF-MM solver implementation following Magron & Févotte (2022).

CRITICAL: H is continuous in [0,1], NOT forced to binary!

Parameters

Y : array-like, shape (m, n)

Binary data matrix {0, 1}

n_components : int

Number of latent components k

max_iter : int

Maximum iterations

tol : float

Convergence tolerance on relative loss change

alpha, beta : float

Beta prior parameters for H

W_init, H_init : array-like, optional

Initial matrices

mask : array-like, optional

Binary mask for observed entries

random_state : int, optional

Random seed

verbose : int

Verbosity level

eps : float

Small constant for numerical stability

Returns

W : array-like, shape (m, k)

Factor matrix with rows on simplex

H : array-like, shape (k, n)

Factor matrix in (0, 1)

losses : list

Loss values per iteration

time_elapsed : float

Total computation time

n_iter : int

Number of iterations run

|||||

Set random seed

if random_state is not None:

np.random.seed(random_state)

Get dimensions

m, n = Y.shape

k = n_components

Initialize mask if not provided

if mask is None:

mask = np.ones_like(Y)

Initialize factors

if W_init is None:

W_init = np.random.uniform(0.1, 0.9, (m, k))

if H_init is None:

H_init = np.random.uniform(0.1, 0.9, (k, n))

CRITICAL: Transpose to match Magron's notation

In Magron's code: W is (k, m) and H is (k, n)

W = W_init.T *# Now (k, m)*

H = H_init.T *# Now (k, n)*

Normalize W columns to sum to 1 (simplex constraint)

W = W / W.sum(axis=0, keepdims=True)

Precompute masked versions for efficiency

```

Y_masked = Y * mask
Y_T = Y.T * mask.T # Transposed masked Y
OneminusY_T = (1 - Y.T) * mask.T # Transposed masked (1-Y)

# Beta prior matrices
A = np.ones_like(H) * (alpha - 1)
B = np.ones_like(H) * (beta - 1)

# Initialize tracking
losses = []
loss_prev = np.inf
start_time = time.time()

# Main optimization loop
for iteration in range(max_iter):

    # ===== H UPDATE =====
    # CRITICAL: H is NOT forced to binary!
    # H stays continuous in (0, 1)

    # Compute  $W^T @ H$ 
    WH = W.T @ H # Shape (m, n)

    # Numerator:  $H * W @ (Y / (WH + \text{eps})) + A$ 
    numerator = H * (W @ (Y_masked / (WH + eps))) + A

    # Denominator:  $(1 - H) * W @ ((1 - Y) / (1 - WH + \text{eps})) + B$ 
    denominator = (1 - H) * (W @ ((1 - Y_masked) / (1 - WH + eps))) + B

    # Update H (stays in (0, 1) naturally)
    H = numerator / (numerator + denominator + eps)

    # Clip to avoid numerical issues at boundaries
    H = np.clip(H, eps, 1 - eps)

    # ===== W UPDATE =====
    # W columns must stay on simplex

    # Compute  $H @ W^T$  (for the transpose computation)
    HW_T = H.T @ W # Shape (n, m)

    # Update W with normalization by n (maintains simplex)
    W_numerator = W * (H @ (Y_T / (HW_T + eps)) + (1 - H) @ (OneminusY_T / (1 - HW_T + eps)))
    W = W_numerator / n

```

```

# Ensure W stays normalized (columns sum to 1)
# This should be maintained by the /n term, but we ensure it for numerical stability
W = W / W.sum(axis=0, keepdims=True)

# ===== COMPUTE LOSS =====
WH = W.T @ H # Recompute after updates

# Log-likelihood term
log_lik = Y_masked * np.log(WH + eps) + (1 - Y_masked) * np.log(1 - WH + eps)

# Prior term
prior = A * np.log(H + eps) + B * np.log(1 - H + eps)

# Total loss (negative log posterior)
loss = -(np.sum(log_lik) + np.sum(prior)) / np.count_nonzero(mask)
losses.append(loss)

if verbose > 0 and iteration % 10 == 0:
    print(f"Iteration {iteration:4d}, Loss: {loss:.6f}")

# ===== CHECK CONVERGENCE =====
if iteration > 0:
    rel_change = abs(loss_prev - loss) / abs(loss_prev)
    if rel_change < tol:
        if verbose > 0:
            print(f"Converged at iteration {iteration} (rel_change: {rel_change:.2e})")
        break

    loss_prev = loss

# Transpose back to our convention
W_final = W.T # Shape (m, k)
H_final = H # Shape (k, n)

time_elapsed = time.time() - start_time
n_iter = iteration + 1

return W_final, H_final, losses, time_elapsed, n_iter

```

3. Integration with NBMF Class

File: `src/nbmf_mm/nbmf.py` (updates needed)


```
from ._solver import nbmf_mm_solver
```

```
class NBMF(BaseEstimator, TransformerMixin):
```

```
    """
```

Non-negative Binary Matrix Factorization via Majorization-Minimization.

IMPORTANT: Despite the name "binary", the factor H is continuous in [0,1] during optimization. The "binary" refers to the input data Y.

```
    """
```

```
def fit(self, X, y=None, mask=None):
```

```
    """Fit NBMF model to binary data X."""
```

```
    # Validate input
```

```
    X = check_array(X, accept_sparse='csr', dtype=np.float64)
```

```
    # Check if data is binary or in [0,1]
```

```
    if not np.all((X >= 0) & (X <= 1)):
```

```
        raise ValueError("X must be in [0,1]")
```

```
    # Handle sparse matrices
```

```
    if sparse.issparse(X):
```

```
        X = X.toarray()
```

```
    # Call the solver with paper-correct implementation
```

```
    W, H, losses, time_elapsed, n_iter = nbmf_mm_solver(
```

```
        Y=X,
```

```
        n_components=self.n_components,
```

```
        max_iter=self.max_iter,
```

```
        tol=self.tol,
```

```
        alpha=self.alpha,
```

```
        beta=self.beta,
```

```
        W_init=self.W_init,
```

```
        H_init=self.H_init,
```

```
        mask=mask,
```

```
        random_state=self.random_state,
```

```
        verbose=self.verbose
```

```
)
```

```
    # Store results
```

```
    self.W_ = W # Shape (n_samples, n_components), rows sum to 1
```

```
    self.components_ = H # Shape (n_components, n_features), values in (0,1)
```

```
    self.loss_curve_ = losses
```

```
self.n_iter_ = n_iter
self.reconstruction_err_ = losses[-1] if losses else np.inf
```

```
return self
```

```
def transform(self, X, mask=None):
```

```
    """Transform X by finding W given fixed H."""
```

```
    check_is_fitted(self, ['components_'])
```

```
    X = check_array(X, accept_sparse='csr', dtype=np.float64)
```

```
    if sparse.issparse(X):
```

```
        X = X.toarray()
```

```
    m = X.shape[0]
```

```
    k = self.n_components
```

```
    H = self.components_
```

```
    # Initialize W randomly
```

```
    W = np.random.uniform(0.1, 0.9, (m, k))
```

```
    # Run a few iterations to find W given fixed H
```

```
    for _ in range(50):
```

```
        # Same W update as in fit, but with fixed H
```

```
        W_T = W.T
```

```
        HW_T = H.T @ W_T
```

```
        if mask is None:
```

```
            Y_T = X.T
```

```
            OneminusY_T = (1 - X).T
```

```
        else:
```

```
            Y_T = X.T * mask.T
```

```
            OneminusY_T = (1 - X).T * mask.T
```

```
        W_T = W_T * (H @ (Y_T / (HW_T + 1e-8)) + (1 - H) @ (OneminusY_T / (1 - HW_T + 1e-8)))
```

```
        W_T = W_T / X.shape[1]
```

```
        W_T = W_T / W_T.sum(axis=0, keepdims=True)
```

```
        W = W_T.T
```

```
    return W
```

```
def inverse_transform(self, W):
```

```
    """Transform W back to data space."""
```

```
    check_is_fitted(self, ['components_'])
```

```
    W = check_array(W, dtype=np.float64)
```



```
# Compute reconstruction  
# Note: W has rows summing to 1, H is in (0,1)  
return W @ self.components_ # Returns probabilities in (0,1)
```

4. Critical Tests to Verify Correctness

File: tests/test_algorithm_correctness.py

```
python
```

```

import numpy as np
import pytest
from nbmf_mm import NBMF

def test_h_continuous_not_binary():
    """Verify H stays continuous, NOT binary during optimization."""
    np.random.seed(42)
    X = (np.random.rand(100, 50) < 0.3).astype(float)

    model = NBMF(n_components=10, max_iter=50)
    model.fit(X)

    H = model.components_

    # H should be continuous in (0, 1), NOT binary
    unique_values = np.unique(H)
    assert len(unique_values) > 2, "H should be continuous, not binary!"
    assert np.all((H >= 0) & (H <= 1)), "H should be in [0, 1]"

    # Check that H has many distinct values (continuous)
    assert len(unique_values) > 100, f"H has only {len(unique_values)} unique values, should be continuous"

    print(f"✓ H is continuous with {len(unique_values)} unique values")

def test_w_simplex_constraint():
    """Verify W rows sum to 1 (simplex constraint)."""
    np.random.seed(42)
    X = (np.random.rand(100, 50) < 0.3).astype(float)

    model = NBMF(n_components=10, max_iter=50)
    model.fit(X)

    W = model.W_
    row_sums = W.sum(axis=1)

    np.testing.assert_allclose(row_sums, 1.0, rtol=1e-5,
                               err_msg="W rows must sum to 1 (simplex constraint)")

    print(f"✓ W rows sum to 1: min={row_sums.min():.6f}, max={row_sums.max():.6f}")

def test_monotonic_convergence():
    """Test that loss decreases monotonically (MM property)."""
    np.random.seed(42)

```

```

X = (np.random.rand(100, 50) < 0.3).astype(float)

model = NBMF(n_components=10, max_iter=100, tol=1e-8)
model.fit(X)

losses = model.loss_curve_

# Check strict monotonicity
violations = []
for i in range(1, len(losses)):
    if losses[i] > losses[i-1] + 1e-12:
        violations.append(i)
        print(f" Violation at iter {i}: {losses[i-1]:.10f} -> {losses[i]:.10f}")

assert len(violations) == 0, f"Found {len(violations)} monotonicity violations!"

print(f"✓ Perfect monotonic convergence over {len(losses)} iterations")

def test_reconstruction_probabilities():
    """Test that reconstruction gives valid probabilities."""
    np.random.seed(42)
    X = (np.random.rand(100, 50) < 0.3).astype(float)

    model = NBMF(n_components=10)
    model.fit(X)

    X_reconstructed = model.inverse_transform(model.W_)

    # Should be probabilities in (0, 1)
    assert np.all((X_reconstructed >= 0) & (X_reconstructed <= 1)), \
        "Reconstructed values should be probabilities in [0, 1]"

    # Should NOT be binary
    unique_recon = np.unique(X_reconstructed)
    assert len(unique_recon) > 100, \
        f"Reconstruction should be continuous probabilities, got {len(unique_recon)} unique values"

    print(f"✓ Reconstruction gives continuous probabilities")

def test_beta_prior_effect():
    """Test that Beta prior parameters affect the solution."""
    np.random.seed(42)
    X = (np.random.rand(50, 30) < 0.3).astype(float)

```

```

# Model with symmetric prior (no preference)
model1 = NBMF(n_components=5, alpha=1.0, beta=1.0, max_iter=100)
model1.fit(X)
H1 = model1.components_

# Model with prior favoring values near 0
model2 = NBMF(n_components=5, alpha=0.5, beta=2.0, max_iter=100, random_state=42)
model2.fit(X)
H2 = model2.components_

# Model with prior favoring values near 1
model3 = NBMF(n_components=5, alpha=2.0, beta=0.5, max_iter=100, random_state=42)
model3.fit(X)
H3 = model3.components_

# Check that priors have expected effect on H
assert H2.mean() < H1.mean(), "Beta(0.5, 2) should push H toward 0"
assert H3.mean() > H1.mean(), "Beta(2, 0.5) should push H toward 1"

print(f"✓ Beta prior affects solution: H means = {H1.mean():.3f}, {H2.mean():.3f}, {H3.mean():.3f}")




```

5. Key Implementation Notes

CRITICAL POINTS:

1. **H is CONTINUOUS in [0,1]**, not binary during optimization
2. **W rows sum to 1** (probability simplex)
3. **Transpose convention:** Internally use Magron's notation (W is $k \times m$, H is $k \times n$), then transpose for sklearn API
4. **Normalization by n:** The W update includes division by n_{features} to maintain simplex
5. **Numerical stability:** Clip H to $[\text{eps}, 1-\text{eps}]$ to avoid $\log(0)$

What We Were Doing Wrong:

-  Forcing H to be binary $\{0,1\}$ after each update
-  Not properly maintaining W simplex constraint
-  Using wrong update equations

What Magron Does Right:

-  H stays continuous in $(0,1)$

- ☒ W normalized to simplex via $/n$ term
- ☒ Proper MM updates that guarantee monotonicity
- ☒ Numerical stability with eps additions

6. Validation Script

Create `examples/validate_magron_implementation.py`:

```
python
```

```
#!/usr/bin/env python3
"""Validate our implementation matches Magron's behavior."""
```

```
import numpy as np
from nbmf_mm import NBMF
import matplotlib.pyplot as plt
```

```
def main():
```

```
    print("="*60)
    print("VALIDATING NBMF-MM IMPLEMENTATION")
    print("="*60)
```

```
    # Generate test data
```

```
    np.random.seed(42)
    X = (np.random.rand(100, 50) < 0.3).astype(float)
    print(f"\nData shape: {X.shape}")
    print(f>Data sparsity: {X.mean():.3f}")
```

```
    # Fit model
```

```
    model = NBMF(
        n_components=10,
        alpha=1.2,
        beta=1.2,
        max_iter=200,
        tol=1e-6,
        verbose=1
    )
    model.fit(X)
```

```
    # Validate results
```

```
    print("\n" + "="*60)
    print("VALIDATION RESULTS")
    print("="*60)
```

```
    W = model.W_
    H = model.components_
    losses = model.loss_curve_
```

```
    # 1. Check H is continuous
```

```
    H_unique = np.unique(H)
    print(f"\n1. H Continuity:")
    print(f"    Unique values in H: {len(H_unique)}")
    print(f"    H range: [{H.min():.4f}, {H.max():.4f}]"
```

```

print(f" H mean: {H.mean():.4f}")
is_continuous = len(H_unique) > 100
print(f" ✓ H is continuous" if is_continuous else " ✗ H is not continuous!")

# 2. Check W simplex constraint
print(f"\n2. W Simplex Constraint:")
row_sums = W.sum(axis=1)
print(f" W row sums: min={row_sums.min():.6f}, max={row_sums.max():.6f}")
simplex_ok = np.allclose(row_sums, 1.0, rtol=1e-5)
print(f" ✓ W rows sum to 1" if simplex_ok else " ✗ W simplex constraint violated!")

# 3. Check monotonic convergence
print(f"\n3. Monotonic Convergence:")
violations = sum(1 for i in range(1, len(losses)) if losses[i] > losses[i-1] + 1e-12)
print(f" Iterations: {len(losses)}")
print(f" Final loss: {losses[-1]:.6f}")
print(f" Monotonicity violations: {violations}")
print(f" ✓ Perfect monotonic convergence" if violations == 0 else f" ✗ {violations} violations!")

# 4. Plot convergence
plt.figure(figsize=(10, 6))
plt.semilogy(losses, 'b-', linewidth=2)
plt.xlabel('Iteration')
plt.ylabel('Loss (log scale)')
plt.title(f'NBMF-MM Convergence (Violations: {violations})')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig('nbmf_convergence_validation.png')
plt.show()

# Overall result
print("\n" + "="*60)
if is_continuous and simplex_ok and violations == 0:
    print("🎉 ALL VALIDATIONS PASSED!")
    print("Implementation correctly follows Magron & Févotte (2022)")
else:
    print("⚠️ SOME VALIDATIONS FAILED")
    print("Check implementation against reference")
print("="*60)

if __name__ == "__main__":
    main()

```

7. Summary of Changes Needed

1. **Update** `_solver.py`: Implement the exact algorithm above with H continuous
2. **Update** `nbfm.py`: Remove any code that forces H to binary
3. **Update tests**: Test for H continuity, not binary constraint
4. **Update documentation**: Clarify that H is continuous, not binary

The key insight is that **H is continuous in $[0,1]$** throughout optimization. The Beta prior naturally encourages H toward 0 or 1 at convergence, making it "effectively binary" for interpretation, but it's never forced to be exactly $\{0,1\}$ during the algorithm!